

OpenLLM, and everything about running LLMs in production

Introduction


Aaron Pham

at aarnphm[_] everywhere

Hobbies: Rock climbing and
reading




BentoML




**Freedom
to build.**

www.bentoml.com

AIGC BOOM!



Emerging LLMs stack (cred. a16z)




Motivation

Running ML in production is... hard

Running LLM in production is even harder!

- Consistency/Hallucination
- Infrastructure Complexity
- Security/Compliance
- Bleeding-edge ecosystem
- Maintainability
- And more




Hosting your own LLMs?

- Customisation/Flexibility
- Security
- Accuracy improvement
- Cost efficiency
- Offline access

BentoML


Introducing, OpenLLM



BentoML

Demo time!

- LLM Inference
- Fine tuning and serve Llama with QLoRA



BentoML

Open-source LLMs

Supports a wide range of architectures and runtime, but not limited to Llama, StableLM, ChatGLM, StarCoder, and more.

Supported models

OpenLLM currently supports the following models. By default, OpenLLM doesn't include dependencies to run all models. The extra model-specific dependencies can be installed with the instructions below.

- ▶ Llama
- ▶ ChatGLM
- ▶ Dolly-v2
- ▶ Falcon
- ▶ Flan-T5
- ▶ GPT-NeoX
- ▶ MPT
- ▶ OPT
- ▶ StableLM
- ▶ StarCoder
- ▶ Baichuan

Supported backends

- PyTorch (Default):

```
openllm start llama --model-id meta-llama/Llama-2-7b-chat-hf --backend pt
```

- vLLM (Recommended):

```
pip install "openllm[llama, vllm]"
openllm start llama --model-id meta-llama/Llama-2-7b-chat-hf --backend vllm
```

BentoML

Built-in optimisation

Token streaming via SSE


```
» curl -N -X 'POST' 'http://44.210.172.220:3000/v1/generate_stream' -H 'accept: application/json' -H 'Content-Type: application/json' -d '{  
  "prompt": "### Instruction:\nWhat is time?\n### Response:",  
  "llm_config": {  
    "use_llama2_prompt": false,  
    "max_new_tokens": 4096,  
    "min_length": 0,  
    "early_stopping": false,  
    "num_beams": 1,  
    "num_beam_groups": 1,  
    "use_cache": true,  
    "temperature": 0.6,  
    "top_k": 12,  
    "top_p": 0.9,  
    "typical_p": 1,  
    "epsilon_cutoff": 0,  
    "eta_cutoff": 0,  
    "diversity_penalty": 0,  
    "repetition_penalty": 1,  
    "encoder_repetition_penalty": 1,  
    "length_penalty": 1,  
    "no_repeat_ngram_size": 0,  
    "renormalize_logits": false,  
    "remove_invalid_values": false,  
    "num_return_sequences": 1,  
    "output_attentions": false,  
    "output_hidden_states": false,  
    "output_scores": false,  
    "encoder_no_repeat_ngram_size": 0,  
    "n": 1,  
    "best_of": 1,  
    "presence_penalty": 0.5,  
    "frequency_penalty": 0,  
    "use_beam_search": false,  
    "ignore_eos": false  
  },  
  "adapter_name": null  
}'
```

BentoML


Built-in optimisation

Continuous batching

- Naive batching
- LLM Inference
- Continuous batching
- Prefill ➔ decode ➔ past values
- PagedAttention
- IO Bound



Given the following prompt:
What is the capital of California:



<https://docs.nvidia.com/deeplearning/performance/dl-performance/benchmarks/back-end/inference/thrust-startup-speedups.html>
<https://huggingface.co/blog/use-dl-inference-to-get-10x-speedup>

Built-in optimisation


Quantisation: GPTQ, kbit

1 MB: GPU Memory required for 1 token of output for 13B model (1 word ~= 1.3 tokens)

512 max new tokens = 512MB VRAM

- Reduce memory footprint
- Improve general throughput

Memory usage for 'NousResearch/Nous-Hermes-Llama2-70b'		
dtype	Largest Layer or Residual Group	Total Size
float32	3.19 GB	256.29 GB
float16/bfloat16	1.6 GB	128.15 GB
int4	408.51 MB	32.04 GB
int8	817.02 MB	64.07 GB



```
1 openllm start llama --model-id TheBloke/Llama-2-13B-chat-GPTQ --quantize gptq --device 0
```

<https://github.com/ray-project/llm-numbers?tab=readme-ov-file#gpu-memory>

Why OpenLLM?

Specialities

- Built-in Inference Optimization with MQA, PagedAttention
- Quantization with GPTQ, k-bit
- Accelerators support with multi-GPUs deployments, TPUs
- Monitoring and evaluation
- Fine-tuning support with qLoRA and various tuning techniques
- Integration with AI tools like LangChain, HuggingFace Agents etc.

Powered by BentoML

- Package model files + dependencies + code into a Bento
- Bentos can be managed and versioned properly in a central place
- Automatically generate docker image for deployment
- Streamlined deployment process: batch inference, online-serving
- Flexible deployment strategy: Docker, Yatai + Kubernetes, bentctl + Terraform, BentoCloud

BentoML

OpenLLM Roadmap

- System prompts
- Unified fine-tuning API for models
- Better CPU inference with GGUF
- Javascript/TypeScript Client library
- OpenAI Compatible APIs
- Optimized modeling for Flash Attentions
- AWQ support, custom CUDA Kernels

Thank you!

Q & A

