

SQL: DATA DEFINITION LANGUAGE



Database Schemas in SQL

2

- SQL is primarily a query language, for getting information from a database.
 - ▣ **Data manipulation language (DML)**
- But SQL also includes a *data-definition* component for describing database schemas.
 - ▣ **Data definition language (DDL)**

Creating (Declaring) a Relation

3

- Simplest form is:

```
CREATE TABLE <name> (  
    <list of elements>  
);
```

- To delete a relation:

```
DROP TABLE <name>;
```

Elements of Table Declarations

4

- Most basic element: an attribute and its type.
- The most common types are:
 - ▣ INT or INTEGER (synonyms).
 - ▣ REAL or FLOAT (synonyms).
 - ▣ CHAR(n) = fixed-length string of n characters.
 - ▣ VARCHAR(n) = variable-length string of up to n characters.

Example: Create Table

5

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price        REAL  
);
```

SQL Values

6

- Integers and reals are represented as you would expect.
- Strings are too, except they require single quotes.
 - ▣ Two single quotes = real quote, e.g., ' Joe' ' s Bar'.
- Any value can be NULL
 - ▣ Unless attribute has NOT NULL constraint
 - ▣ E.g., price REAL not null,

Dates and Times

7

- DATE and TIME are types in SQL.

- The form of a date value is:

DATE 'yyyy-mm-dd'

- ▣ **Example:** DATE '2007-09-30' for Sept. 30, 2007.

Times as Values

8

- The form of a time value is:

TIME 'hh:mm:ss'

with an optional decimal point and fractions of a second following.

- **Example:** TIME '15:30:02.5' = two and a half seconds after 3:30PM.

Declaring Keys

9

- An attribute or list of attributes may be declared **PRIMARY KEY** or **UNIQUE**.
- Either says that no two tuples of the relation may agree in all the attribute(s) on the list.

Our Running Example

10

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)

- Underline = **key** (tuples cannot have the same value in all key attributes).

Declaring Single-Attribute Keys

11

- Place PRIMARY KEY or UNIQUE after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE Beers (  
    name    CHAR(20)  UNIQUE,  
    manf    CHAR(20)  
);
```

Declaring Multiattribute Keys

12

- A key declaration can also be another element in the list of elements of a CREATE TABLE statement.
- This form is essential if the key consists of more than one attribute.
 - ▣ May be used even for one-attribute keys.

Example: Multiattribute Key

13

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price        REAL ,  
    PRIMARY KEY (bar, beer)  
);
```

PRIMARY KEY vs. UNIQUE

14

1. There can be only one PRIMARY KEY for a relation, but several UNIQUE attributes.
2. No attribute of a PRIMARY KEY can ever be NULL in any tuple. But attributes declared UNIQUE may have NULL's, and there may be several tuples with NULL.

Kinds of Constraints

15

- Keys
- Foreign-key, or referential-integrity.
- Domain constraints
 - ▣ Constrain values of a particular attribute.
- Tuple-based constraints
 - ▣ Relationship among components.
- Assertions: any SQL boolean expression

Foreign Keys

16

- Values appearing in attributes of one relation must appear together in certain attributes of another relation.
- **Example:** in `Sells(bar, beer, price)`, we might expect that a beer value also appears in `Beers.name`

Expressing Foreign Keys

17

- Use keyword REFERENCES, either:
 1. After an attribute (for one-attribute keys).
 2. As an element of the schema:

FOREIGN KEY (<list of attributes>)

REFERENCES <relation> (<attributes>)

- Referenced attributes must be declared PRIMARY KEY or UNIQUE.

Example: With Attribute

18

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar       CHAR(20),  
    beer      CHAR(20) REFERENCES Beers(name),  
    price     REAL );
```

Example: As Schema Element

19

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar        CHAR(20),  
    beer       CHAR(20),  
    price      REAL,  
    FOREIGN KEY (beer) REFERENCES  
        Beers (name) );
```

Enforcing Foreign-Key Constraints

20

- If there is a foreign-key constraint from relation R to relation S , two violations are possible:
 1. An insert or update to R introduces values not found in S .
 2. A deletion or update to S causes some tuples of R to “dangle.”

Actions Taken --- (1)

21

- **Example:** suppose $R = \text{Sells}$, $S = \text{Beers}$.
- An insert or update to **Sells** that introduces a nonexistent beer must be rejected.
- A deletion or update to **Beers** that removes a beer value found in some tuples of **Sells** can be handled in three ways...

Actions Taken --- (2)

22

1. *Default* : Reject the modification.
2. *Cascade* : Make the same changes in Sells.
 - ▣ Deleted beer: delete Sells tuple.
 - ▣ Updated beer: change value in Sells.
3. *Set NULL* : Change the beer to NULL.

Example: Cascade

23

- Delete the Bud tuple from Beers:
 - ▣ Then delete all tuples from Sells that have beer = 'Bud'.
- Update the Bud tuple by changing 'Bud' to 'Budweiser':
 - ▣ Then change all Sells tuples with beer = 'Bud' to beer = 'Budweiser'.

Example: Set NULL

24

- Delete the Bud tuple from Beers:
 - ▣ Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- Update the Bud tuple by changing 'Bud' to 'Budweiser':
 - ▣ Same change as for deletion.

Choosing a Policy

25

- When we declare a foreign key, we may choose policies SET NULL or CASCADE independently for deletions and updates.
- Follow the foreign-key declaration by:
ON [UPDATE, DELETE][SET NULL CASCADE]
- Two such clauses may be used.
- Otherwise, the default (reject) is used.

Example: Setting Policy

26

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    FOREIGN KEY (beer)  
        REFERENCES Beers (name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```

Attribute-Based Checks

27

- Constraints on the value of a particular attribute.
- Add `CHECK(<condition>)` to the declaration for the attribute.
- The condition may use the name of the attribute, but any other relation or attribute name must be in a subquery.

Example: Attribute-Based Check

28

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20) CHECK ( beer IN  
        (SELECT name FROM Beers) ),  
    price    REAL CHECK ( price <= 5.00 )  
);
```

Timing of Checks

29

- Attribute-based checks are performed only when a value for that attribute is inserted or updated.
 - ▣ **Example:** `CHECK (price <= 5.00)` checks every new price and rejects the modification (for that tuple) if the price is more than \$5.
 - ▣ **Example:** `CHECK (beer IN (SELECT name FROM Beers))` not checked if a beer is deleted from Beers (unlike foreign-keys).

Tuple-Based Checks

30

- CHECK (<condition>) may be added as a relation-schema element.
- The condition may refer to any attribute of the relation.
 - ▣ But other attributes or relations require a subquery.
- Checked on insert or update only.

Example: Tuple-Based Check


31

- Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         CHAR(20) ,  
    price        REAL,  
    CHECK (bar = 'Joe''s Bar' OR  
           price <= 5.00)  
);
```

Asg 1 Update: Missing attribute in Baggage

1



INTRODUCTION TO SQL



Why SQL?

3

- SQL is a very-high-level language.
 - ▣ Structured Query Language
 - ▣ Say “what to do” rather than “how to do it.”
 - ▣ Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query.
 - ▣ Called “query optimization.”

Database Schemas in SQL

4

- SQL is primarily a query language, for getting information from a database.
 - ▣ **Data manipulation language (DML)**
- But SQL also includes a *data-definition* component for describing database schemas.
 - ▣ **Data definition language (DDL)**

Select-From-Where Statements

5

SELECT desired attributes

FROM one or more tables

WHERE condition about tuples of
the tables

Our Running Example

6

- Our SQL queries will be based on the following database schema.
 - ▣ Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)

Example

7

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch';
```

Result of Query

8

name

Bud

Bud Lite

Michelob

...

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

Meaning of Single-Relation Query

9

- Begin with the relation in the FROM clause.
- Apply the selection indicated by the WHERE clause.
- Apply the extended projection indicated by the SELECT clause.

Operational Semantics - General

10

- Think of a *tuple variable* visiting each tuple of the relation mentioned in FROM.
- Check if the tuple assigned to the tuple variable satisfies the WHERE clause.
- If so, compute the attributes or expressions of the SELECT clause using the components of this tuple.

Operational Semantics

11

name	manf
Bud	Anheuser-Busch

If so, include t.name
in the result

Check if
Anheuser-Busch

Tuple-variable t
loops over all
tuples

Example

12

- What beers are made by Anheuser-Busch?

```
SELECT name  
FROM Beers  
WHERE manf = 'Anheuser-Busch';
```

OR:

```
SELECT t.name  
FROM Beers t  
WHERE t.manf = 'Anheuser-Busch';
```

Note: these are identical queries.

* In SELECT clauses

13

- When there is one relation in the FROM clause, * in the SELECT clause stands for “all attributes of this relation.”
- **Example:** Using **Beers(name, manf):**

```
SELECT *
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch';
```

Result of Query:

14

name	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
...	...

Now, the result has each of the attributes of Beers.

Renaming Attributes

15

- If you want the result to have different attribute names, use “AS <new name>” to rename an attribute.

- **Example:** Using **Beers(name, manf):**

```
SELECT name AS beer, manf
FROM Beers
WHERE manf = 'Anheuser-Busch'
```

Result of Query:

16

beer	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
...	...

Expressions in SELECT Clauses

17

- Any valid expression can appear as an element of a SELECT clause.

- **Example:** Using `Sells(bar, beer, price)`:

```
SELECT bar, beer,  
       price*95 AS priceInYen  
FROM Sells;
```


Result of Query

18

bar	beer	priceInYen
Joe's	Bud	285
Sue's	Miller	342
...

Example: Constants as Expressions

19

- Using `Likes(drinker, beer)`:

```
SELECT drinker,  
       'likes Bud' AS whoLikesBud  
FROM Likes  
WHERE beer = 'Bud';
```

Result of Query

20

drinker	whoLikesBud
Sally	likes Bud
Fred	likes Bud
...	...

Complex Conditions in WHERE Clause

21

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.

Example: Complex Condition

22

- Using `Sells(bar, beer, price)`, find the price Joe's Bar charges for Bud:

```
SELECT  price
FROM    Sells
WHERE   bar = 'Joe''s Bar' AND
        beer = 'Bud';
```

Patterns

23

- A condition can compare a string to a pattern by:
 - ▣ `<Attribute> LIKE <pattern>` or `<Attribute> NOT LIKE <pattern>`
- *Pattern* is a quoted string
 - ▣ `%` = “any string”;
 - ▣ `_` = “any character”.

Example: LIKE

24

- Using `Drinkers(name, addr, phone)` find the drinkers with exchange 555:

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555-__ __ __';
```

NULL Values

25

- Tuples in SQL relations can have NULL as a value for one or more components.
- Meaning depends on context. Two common cases:
 - ▣ *Missing value* : e.g., we know Joe's Bar has some address, but we don't know what it is.
 - ▣ *Inapplicable* : e.g., the value of attribute *spouse* for an unmarried person.

Comparing NULL's to Values

26

- The logic of conditions in SQL is really 3-valued logic: TRUE, FALSE, UNKNOWN.
- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer iff the WHERE clause is TRUE (not FALSE or UNKNOWN).

Three-Valued Logic

27

- To understand how AND, OR, and NOT work in 3-valued logic
- For TRUE result
 - ▣ OR: at least one operand must be TRUE
 - ▣ AND: both operands must be TRUE
 - ▣ NOT: operand must be FALSE
- For FALSE result
 - ▣ OR: both operands must be FALSE
 - ▣ AND: at least one operand must be FALSE
 - ▣ NOT: operand must be TRUE
- Otherwise, result is UNKNOWN

Example

28


- From the following Sells relation:

bar	beer	price
Joe's Bar	Bud	NULL

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 5.00;



Multi-Relation Queries

29

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>” .

Example: Joining Two Relations

30

- Using relations `Likes(drinker, beer)` and `Frequents(drinker, bar)`, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes, Frequents
WHERE bar = 'Joe''s Bar' AND
       Frequents.drinker = Likes.drinker;
```

Example: Joining Two Relations

31

- Alternatively can use explicit (named) tuple variables

```
SELECT beer
FROM Likes l, Frequents f
WHERE bar = 'Joe''s Bar' AND
       f.drinker = l.drinker;
```

Formal Semantics

32

- Almost the same as for single-relation queries:
 - Start with the product of all the relations in the FROM clause.
 - Apply the selection condition from the WHERE clause.
 - Project onto the list of attributes and expressions in the SELECT clause.

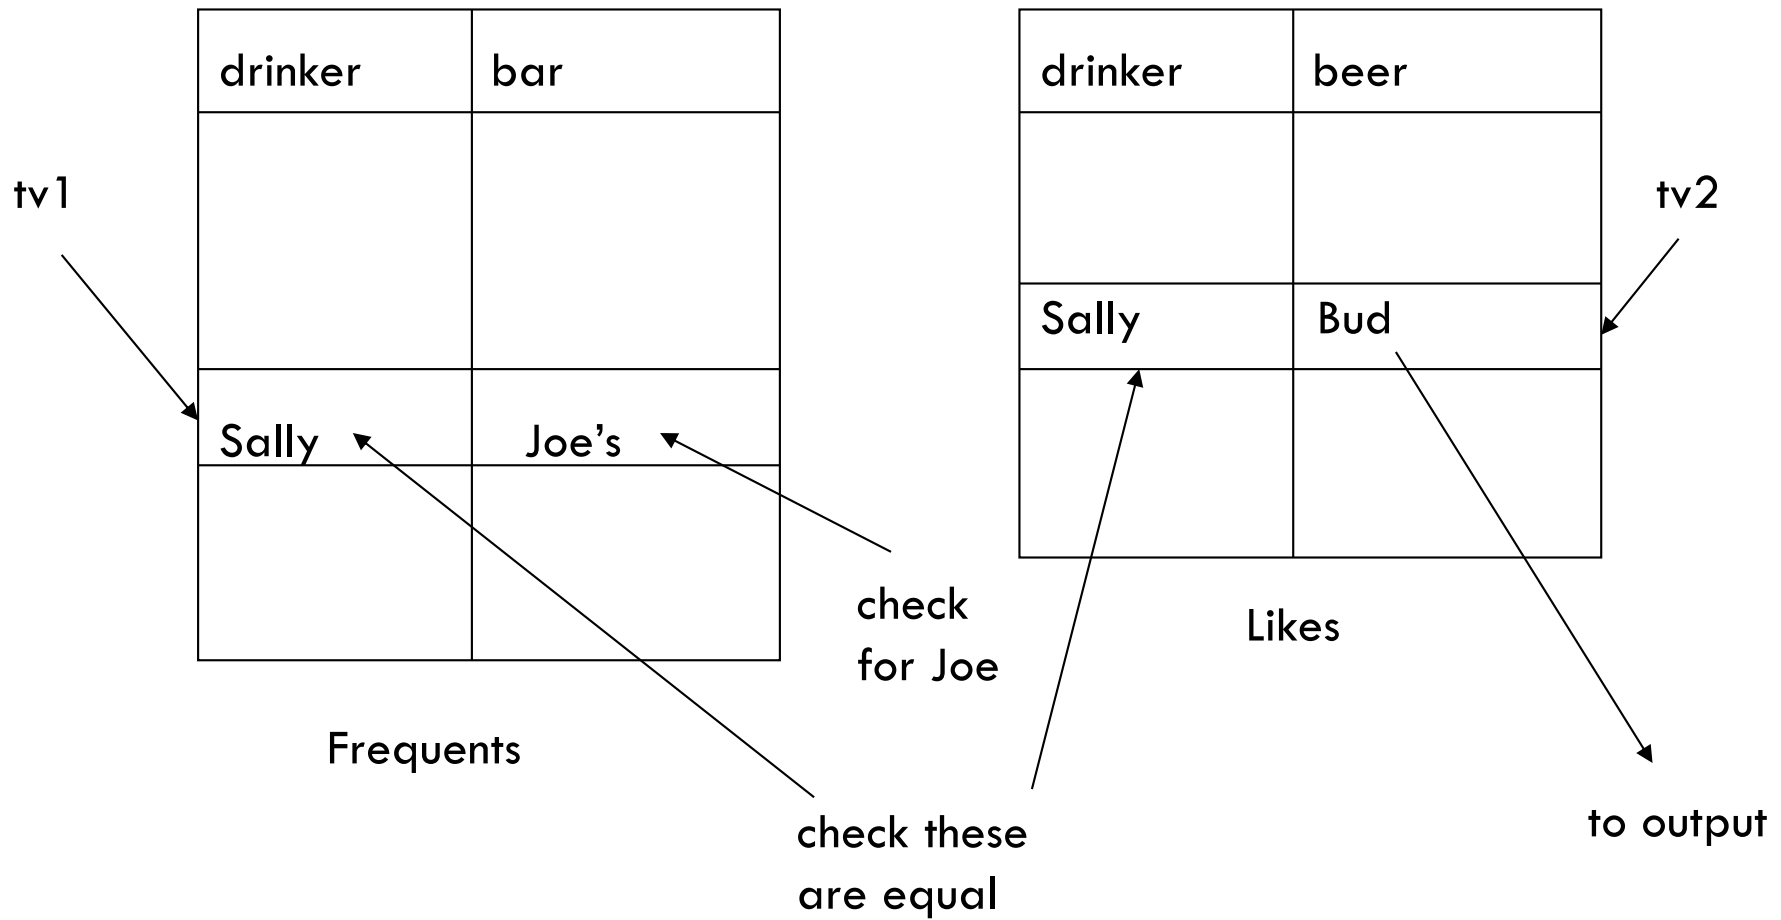
Operational Semantics

33

- Imagine one tuple-variable for each relation in the FROM clause.
 - ▣ These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.

Example

34



Explicit Tuple-Variables

35

- Sometimes, a query needs to use two copies of the same relation.
- Distinguish copies by following the relation name by the name of a tuple-variable, in the FROM clause.
- It's always an option to rename relations this way, even when not essential.

Example: Self-Join

36

- From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.
 - ▣ Do not produce pairs like (Bud, Bud).
 - ▣ Do not produce the same pair twice like (Bud, Miller) and (Miller, Bud).

Select-From-Where Statements

1

SELECT desired attributes

FROM one or more tables

WHERE condition about tuples of
the tables

Example

2

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch';
```

Result of Query

3

name

Bud

Bud Lite

Michelob

...

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

Operational Semantics

4

name	manf
Bud	Anheuser-Busch

Tuple-variable t
loops over all
tuples

Check if
Anheuser-Busch

If so, include $t.name$
in the result

Comparing NULL's to Values

5

- The logic of conditions in SQL is really 3-valued logic: TRUE, FALSE, UNKNOWN.
- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer iff the WHERE clause is TRUE (not FALSE or UNKNOWN).

Three-Valued Logic

6

- To understand how AND, OR, and NOT work in 3-valued logic
- For TRUE result
 - ▣ OR: at least one operand must be TRUE
 - ▣ AND: both operands must be TRUE
 - ▣ NOT: operand must be FALSE
- For FALSE result
 - ▣ OR: both operands must be FALSE
 - ▣ AND: at least one operand must be FALSE
 - ▣ NOT: operand must be TRUE
- Otherwise, result is UNKNOWN

Example

7


- From the following Sells relation:

bar	beer	price
Joe's Bar	Bud	NULL

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 5.00;



Multi-Relation Queries

8

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>” .

Example: Joining Two Relations

9

- Using relations `Likes(drinker, beer)` and `Frequents(drinker, bar)`, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes, Frequents
WHERE bar = 'Joe''s Bar' AND
       Frequents.drinker = Likes.drinker;
```

Example: Joining Two Relations

10

- Alternatively can use explicit (named) tuple variables

```
SELECT beer
FROM Likes l, Frequents f
WHERE bar = 'Joe''s Bar' AND
       f.drinker = l.drinker;
```

Formal Semantics

11

- Almost the same as for single-relation queries:
 - Start with the product of all the relations in the FROM clause.
 - Apply the selection condition from the WHERE clause.
 - Project onto the list of attributes and expressions in the SELECT clause.

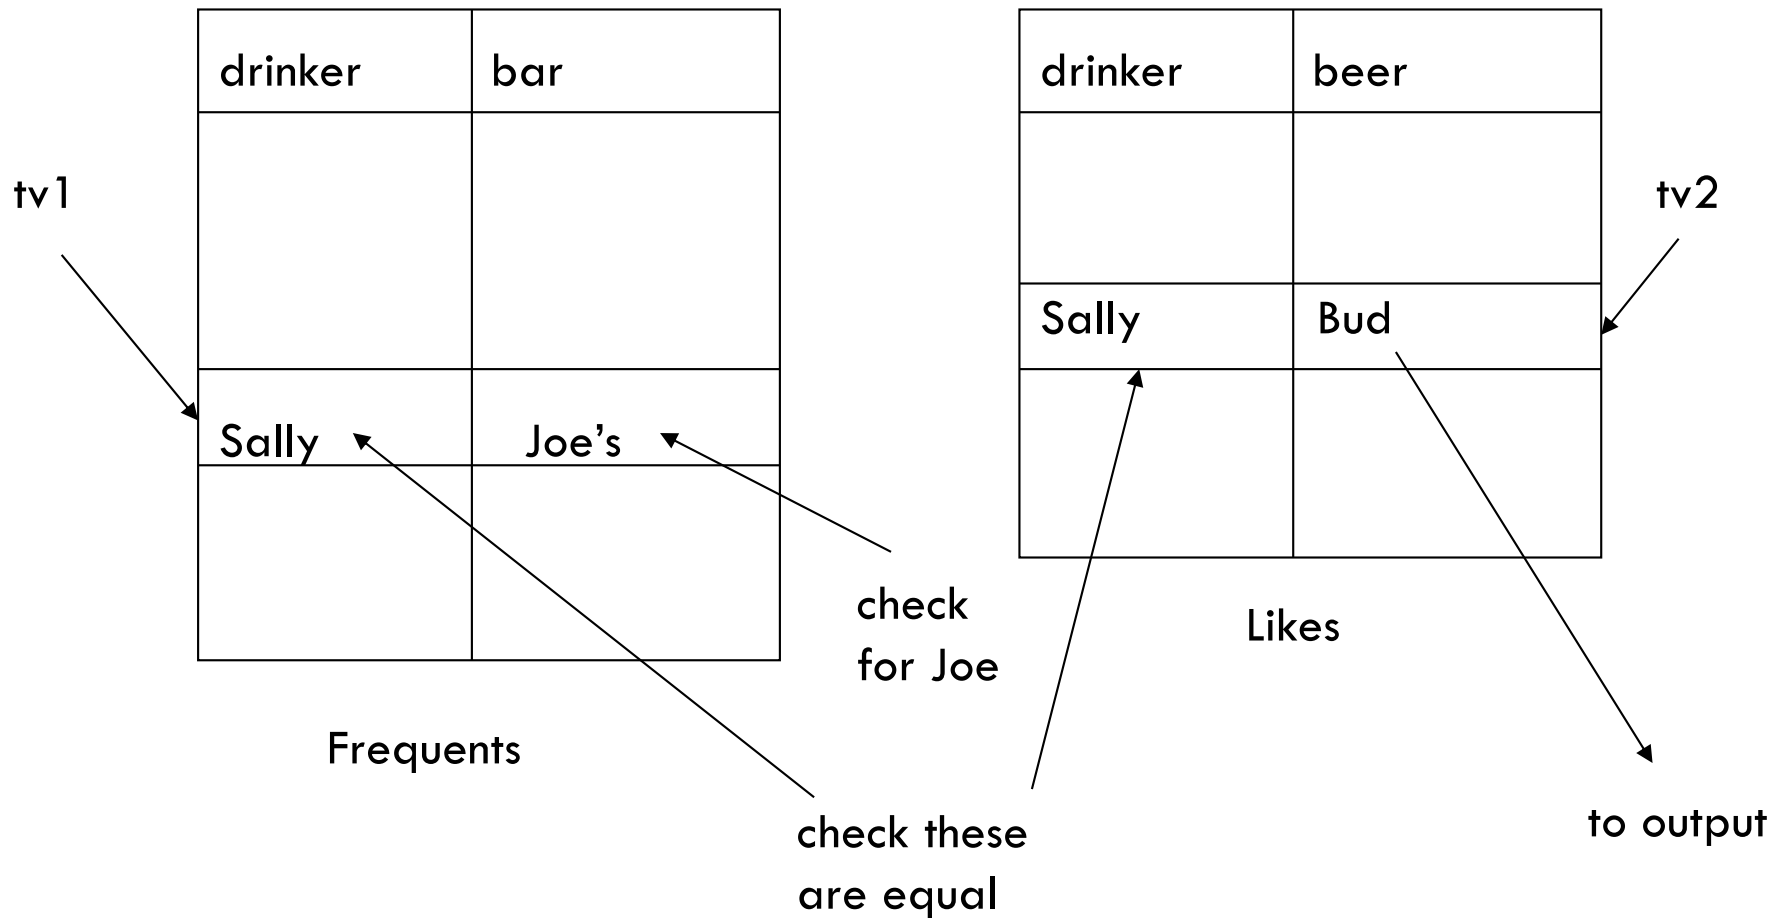
Operational Semantics

12

- Imagine one tuple-variable for each relation in the FROM clause.
 - ▣ These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.

Example

13



Explicit Tuple-Variables

14

- Sometimes, a query needs to use two copies of the same relation.
- Distinguish copies by following the relation name by the name of a tuple-variable, in the FROM clause.
- It's always an option to rename relations this way, even when not essential.

Example: Self-Join

15

- From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.
 - ▣ Do not produce pairs like (Bud, Bud).
 - ▣ Do not produce the same pair twice like (Bud, Miller) and (Miller, Bud).

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND
      b1.name < b2.name;
```

Subqueries

16

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including FROM and WHERE clauses.
- **Example:** in place of a relation in the FROM clause, we can use a subquery and then query its result.
 - Must use a tuple-variable to name tuples of the result.

Example: Subquery in FROM

17

- Find the beers liked by at least one person who frequents Joe's Bar.

Drinkers who
frequent Joe's Bar

```
SELECT beer
```

```
FROM Likes, (SELECT drinker  
              FROM Frequents  
              WHERE bar = 'Joe's Bar') JD
```

```
WHERE Likes.drinker = JD.drinker;
```

Subqueries often obscure queries

18

- Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes l, Frequents f
WHERE l.drinker = f.drinker AND
      bar = 'Joe's Bar';
```

Simple join query

Subqueries That Return One Tuple

19

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value.
 - ▣ Usually, the tuple has one component.
 - ▣ Remember SQL's 3-valued logic.

Example: Single-Tuple Subquery

20

- Using `Sells(bar, beer, price)`, find the bars that serve Miller for the same price Joe charges for Bud.

Two queries would work:

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.

Query + Subquery Solution

21

SELECT bar

FROM Sells


WHERE beer = 'Miller' AND price

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.

Sells(bar, beer, price)

= (SELECT price
FROM Sells
WHERE bar = 'Joe''s Bar'
AND beer = 'Bud');

The price at
which Joe
sells Bud



What if price of Bud is NULL?


Query + Subquery Solution

22

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND
```

```
price = (SELECT price  
        FROM Sells  
        WHERE beer = 'Bud');
```

What if subquery
returns multiple
values?



Recap: Conditions in WHERE Clause

23

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
- LIKE operator

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, $\geq \$90,000$ and $\leq \$100,000$)
 - **select** *name*
 from *instructor*
 where *salary* **between** 90000 **and** 100000

Subqueries

36

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including FROM and WHERE clauses.
- **Example:** in place of a relation in the FROM clause, we can use a subquery and then query its result.
 - ▣ Must use a tuple-variable to name tuples of the result.

Example: Subquery in FROM

37

- Find the beers liked by at least one person who frequents Joe's Bar.

Drinkers who
frequent Joe's Bar

```
SELECT beer
```

```
FROM Likes, (SELECT drinker  
              FROM Frequents  
              WHERE bar = 'Joe's Bar') JD
```

```
WHERE Likes.drinker = JD.drinker;
```

Subqueries often obscure queries

38

- Find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes l, Frequents f
WHERE l.drinker = f.drinker AND
      bar = 'Joe's Bar';
```

Simple join query

Subqueries That Return One Tuple

39

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value.
 - ▣ Usually, the tuple has one component.
 - ▣ Remember SQL's 3-valued logic.

Example: Single-Tuple Subquery

40

- Using `Sells(bar, beer, price)`, find the bars that serve Miller for the same price Joe charges for Bud.

Two queries would work:

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.

Query + Subquery Solution

41

SELECT bar

FROM Sells


WHERE beer = 'Miller' AND price

- Find the price Joe charges for Bud.
- Find the bars that serve Miller at that price.

Sells(bar, beer, price)

= (SELECT price
FROM Sells
WHERE bar = 'Joe''s Bar'
AND beer = 'Bud');

The price at
which Joe
sells Bud



What if price of Bud is NULL?


Query + Subquery Solution

42

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND
```

```
price = (SELECT price  
        FROM Sells  
        WHERE beer = 'Bud');
```

What if subquery
returns multiple
values?



Recap: Conditions in WHERE Clause

43

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
- LIKE operator

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, \geq \$90,000 and \leq \$100,000)
 - **select** *name*
from *instructor*
where *salary* **between** 90000 **and** 100000

The Operator ANY

44

- $x = \text{ANY}(<\text{subquery}>)$ is a boolean condition that is true iff x equals at least one tuple in the subquery result.
 - ▣ $=$ could be any comparison operator.
- **Example:** $x \geq \text{ANY}(<\text{subquery}>)$ means x is not the uniquely smallest tuple produced by the subquery.
 - ▣ Note tuples must have one component only.

The Operator ALL

45

- $x \langle \rangle \text{ALL}(\langle \text{subquery} \rangle)$ is true iff for every tuple t in the relation, x is not equal to t .
 - ▣ That is, x is not in the subquery result.
- $\langle \rangle$ can be any comparison operator.
- **Example:** $x \geq \text{ALL}(\langle \text{subquery} \rangle)$ means there is no tuple larger than x in the subquery result.

Example: ALL

46

- From **Sells(bar, beer, price)**, find the beer(s) sold for the highest price.

SELECT beer

FROM Sells

WHERE price \geq

ALL(SELECT price
FROM Sells);

price from the outer
Sells must not be
less than any price.

The IN Operator

47

- `<value> IN (<subquery>)` is true if and only if the `<value>` is a member of the relation produced by the subquery.
 - ▣ Opposite: `<value> NOT IN (<subquery>)`.
- IN-expressions can appear in WHERE clauses.
- `WHERE col IN (value1, value2, ...)`

IN is Concise

48

- ❑ `SELECT * FROM Cartoons`
`WHERE LastName IN ('Simpsons', 'Smurfs', 'Flintstones')`

- ❑ `SELECT * FROM Cartoons`
`WHERE LastName = 'Simpsons'`
`OR LastName = 'Smurfs'`
`OR LastName = 'Flintstones'`

Example: IN

49

- Using **Beers(name, manf)** and **Likes(drinker, beer)**, find the name and manufacturer of each beer that Fred likes.

```
SELECT *
```

```
FROM Beers
```

```
WHERE name IN (SELECT beer
```

The set of
beers Fred
likes

```
FROM Likes
```

```
WHERE drinker = 'Fred');
```


IN vs. Join

50

```
SELECT R.a  
FROM R, S  
WHERE R.b = S.b;
```

```
SELECT R.a  
FROM R  
WHERE b IN (SELECT b FROM S);
```

IN is a Predicate About R's Tuples

51

```
SELECT a
FROM R
WHERE b IN
```

```
(SELECT b FROM S);
```

Two 2's

a	b
1	2
3	4

R

b	c
2	5
2	6

S

(1,2) satisfies
the condition;
1 is output once.

One loop, over
the tuples of R

This Query Pairs Tuples from R, S

52

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

a	b
1	2
3	4

R

b	c
2	5
2	6

S

Double loop, over
the tuples of R and S

(1,2) with (2,5)
and (1,2) with
(2,6) both satisfy
the condition;
1 is output twice.


Query + Subquery Solution

42

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND
```

```
price = (SELECT price  
        FROM Sells  
        WHERE beer = 'Bud');
```

What if subquery
returns multiple
values?



The Operator ANY

43

- $x = \text{ANY}(<\text{subquery}>)$ is a boolean condition that is true iff x equals at least one tuple in the subquery result.
 - ▣ $=$ could be any comparison operator.
- **Example:** $x \geq \text{ANY}(<\text{subquery}>)$ means x is not the uniquely smallest tuple produced by the subquery.
 - ▣ Note tuples must have one component only.

The Operator ALL

44

- $x \langle \rangle \text{ALL}(\langle \text{subquery} \rangle)$ is true iff for every tuple t in the relation, x is not equal to t .
 - ▣ That is, x is not in the subquery result.
- $\langle \rangle$ can be any comparison operator.
- **Example:** $x \geq \text{ALL}(\langle \text{subquery} \rangle)$ means there is no tuple larger than x in the subquery result.

The IN Operator

45

- `<value> IN (<subquery>)` is true if and only if the `<value>` is a member of the relation produced by the subquery.
 - ▣ Opposite: `<value> NOT IN (<subquery>)`.
- IN-expressions can appear in WHERE clauses.
- `WHERE col IN (value1, value2, ...)`

IN vs. Join

46

```
SELECT R.a  
FROM R, S  
WHERE R.b = S.b;
```

```
SELECT R.a  
FROM R  
WHERE b IN (SELECT b FROM S);
```


IN is a Predicate About R's Tuples

47

```
SELECT a
FROM R
WHERE b IN
```

```
(SELECT b FROM S);
```

Two 2's

a	b
1	2
3	4

R

b	c
2	5
2	6

S

(1,2) satisfies
the condition;
1 is output once.

One loop, over
the tuples of R

This Query Pairs Tuples from R, S

48

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

a	b
1	2
3	4

R

b	c
2	5
2	6

S

Double loop, over
the tuples of R and S

(1,2) with (2,5)
and (1,2) with
(2,6) both satisfy
the condition;
1 is output twice.

Back to our original query...

49

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND  
price = (SELECT price
```

```
FROM Sells  
WHERE beer = 'Bud');
```

Use IN() or = ANY()

Recap

50

- `IN()` is equivalent to `= ANY()`
- For `ANY()`, you can use other comparison operators such as `>`, `<`,... etc, but not applicable for `IN()`
- The `< >ANY` operator, however, differs from `NOT IN`:
 - ▣ `< >ANY` means `not = a, or not = b, or not = c`
 - ▣ `NOT IN` means `not = a, and not = b, and not = c.`
 - ▣ `<>ALL` means the same as `NOT IN`.

Example: =ANY

51

Sells

Bar	Beer	Price
Jane	Miller	3.00
Joe	Miller	4.00
Joe	Bud	3.00
Jack	Bud	4.00
Tom	Miller	4.50

```
SELECT Bar
FROM Sells
WHERE Beer = 'Miller' AND Price =
      ANY(SELECT Price
           FROM Sells
           WHERE Beer='Bud')
```

Result

Bar
Jane
Joe

The Exists Operator

52

- EXISTS(<subquery>) is true if and only if the subquery result is not empty.
- **Example:** From **Beers(name, manf)** , find those beers that are the unique (only) beer made by their manufacturer.

Example: EXISTS

53

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS (
```

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute. (Some DBMS consider this ambiguous.)

Set of beers with the same manf as b1, but not the same beer

```
SELECT *  
FROM Beers  
WHERE manf = b1.manf AND  
      name <> b1.name);
```

Notice the SQL “not equals” operator


Union, Intersection, and Difference

54

- Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:
 - ▣ (<subquery>) UNION (<subquery>)
 - ▣ (<subquery>) INTERSECT (<subquery>)
 - ▣ (<subquery>) EXCEPT (<subquery>)

Visually

55



Example: Intersection

56

- Using `Likes(drinker, beer)`, `Sells(bar, beer, price)`, and `Frequents(drinker, bar)`, find the drinkers and beers such that:
 - The drinker likes the beer, and
 - The drinker frequents at least one bar that sells the beer.

Solution

subquery is
really a stored
table.

```
(SELECT * FROM Likes)
```

INTERSECT

```
(SELECT drinker, beer  
FROM Sells, Frequents  
WHERE Frequents.bar = Sells.bar  
);
```

The drinker frequents
a bar that sells the
beer.